

Machine Learning

Gcinizwe Dlamini




Agenda

- A quick recap of the last lecture
- PyTorch Components
 - Tensors
 - Autograd
 - Optimizers & Loss functions
 - Model
- PyTorch model training
- Saving model and Deployment
- Hands-on example
- Q & A

Recap

Data Augmentation



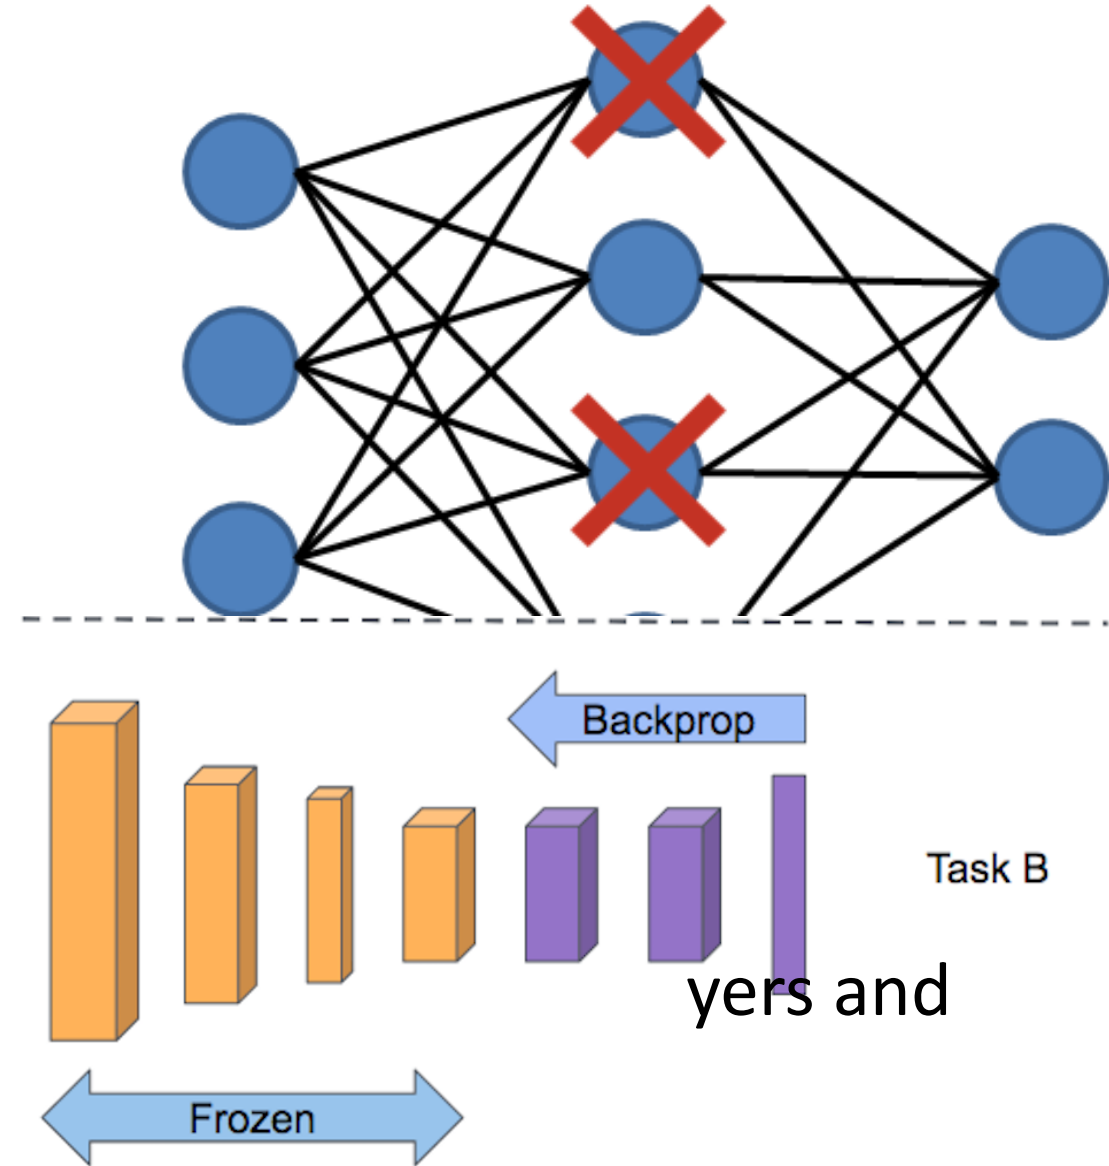
Why Deep Learning? Slide by Andrew Ng

- How machine learning scale with amount of data
- So let's try to increase our DNN
- **NB:** Creating duplicate data leads to overfitting
- Data augmentation is a simple way that helps to avoid overfitting

Transfer Learning

Input A

Input B



rarely used after convolutional layers


Recap (2)

Early Stopping

- While training large capacity n loss will be steadily decreasing kind of U-shape curve decreases
- Stop the training at the lowest validation loss point

[Must read Resource](#)

Loss (negative log-likelihood)



Batch normalization (2)

- Standardizes the input mini-batch before being fed to the next layer
- For inferencing **Exponential Moving Average** of the mean and variance are calculated
- ! and " are learned parameters

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots x_m\}$;
Parameters to be learned: γ, β
Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$
$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$
$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$
$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

Algorithm 1: Batch Normalizing Transform, applied to activation x over a mini-batch.

Tensors

- Tensors are similar to NumPy's ndarrays
- Can run on GPUs or other hardware accelerators
- Optimized for automatic differentiation

Autograd

- Automatic differentiation package : No need to worry about back propagation partial derivatives and chain rule
- Tensors track their computational history and support gradient computation
- **requires_grad=True** : Tells PyTorch that we want to compute gradients for the specific tensor

Autograd : **backward()**

- The **backward()** function is responsible for calculation of gradients and **accumulate** (not apply) them in respective tensors
- The tensor with **requires_grad=True**: has attribute to check the gradients values : '**grad**'
- Because of the **accumulate** it is important to zero the accumulated values before any calculations '**zero_()**'

Optimizers

- Optimizers facilitate the update of tensors values with the gradients
- In PyTorch the reset of accumulated gradients is facilitated by the optimizer '**zero_grad()**'
- To facilitate the updates of tensors values with the gradient values and learning rate '**step()**' method should be evoked

Loss Functions

Recap on Linear Regression Loss function

Mean Squared Error (MSE)

$$f(x_i) = w_0 + w_1 x_i$$

$$e_i = y_i - f(x_i)$$

$$\mathcal{L}(w_0, w_1) = \frac{1}{n} \sum_{i=1}^n (y_i - (w_0 + w_1 x_i))^2$$

We need to find the value of parameters that minimize this cost or loss function.

`torch.nn.MSELoss`

Recap on Logistic Regression Loss Function

Loss Function of Logistic Regression (4)

$$\mathcal{L}(\mathbf{w}) = -\frac{1}{n} \sum_{i=1}^n y^i \log(p(x^i)) + (1 - y^i) \log(1 - p(x^i))$$

- Given this loss function, what do you think we are gonna do next?
- We will define our objective function

$$\underset{w_0, w_1}{\operatorname{argmin}} \mathcal{L}(\mathbf{w})$$

`torch.nn.BCELoss`

Other Loss functions

- Kullback-Leibler divergence : **`torch.nn.KLDivLoss`**
- Cosine Embedding : **`torch.nn.CosineEmbeddingLoss`**
- Negative log likelihood loss : **`torch.nn.NLLLoss`**
- Cross entropy loss : **`nn.CrossEntropyLoss`**

In PyTorch a loss function is called ***crit**erion*

[More Loss functions and description](#)

Creating a simple ANN & DNN

Model

- A **model** is represented by a regular **Python class** that inherits from the **Module** class
- **`__init__(self)`**: it defines the parts that make up the model
- **`forward(self, x)`**: performs a forward pass

```
import torch

class myModel(torch.nn.Module):
    def __init__(self) :
        super(...).__init__()
        ...
    def forward(self,x):
        ...
        return x
```



Model important attributes

- **`model.train()`** : sets the model to training mode. Keep track of the gradients and computations in the graph
- **`model.eval()`** : sets the model to evaluation mode (no need to accumulate gradients and ignore dropout)
- **`model.parameters()`** : retrieves an iterator over all model's parameters
- **`model.state_dict()`** : retrieves model current values for all parameters

Sequential Models

```
import torch

class myModel(torch.nn.Module):
    def __init__(self) :
        super(...).__init__()
        ...
    def forward(self,x):
        ...
```



```
import torch

model = torch.nn.Sequential(...)
model.train()
...
model.eval()
```

Training the model

Typical procedure in training a neural network using PyTorch:

1. Define the model
2. Define loss function
3. Define optimizer
4. Define training loop

```
import torch
import torch.optim as optim

# 1. define model
model = torch.nn.Sequential(...)

# 2. define loss function (i.e regression)
loss_function = torch.nn.MSELoss()

# 3. Define Optimizer
optimizer = optim.SGD(model.parameters(),
                      lr=0.1)

# define training loop

Next slide ...
```

Training the model

Typical procedure in training a neural network using PyTorch:

1. Define the model
2. Define loss function
3. Define optimizer
- 4. Define training loop**

```
import torch

# 4. Define training loop
for epoch in range(n_epochs):
    for batch in batches:
        inputs = batch[0].to(device)
        labels = batch[1].to(device)

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = net(inputs)
        loss = loss_function(outputs, labels)

        # accumulate gradients and update params
        loss.backward()
        optimizer.step()
```

Saving and loading Model

- PyTorch models store the learned parameters in an internal state dictionary `model.state_dict()`
- Can be persisted via the PyTorch save method `torch.save(...)`

```
import torch
model = torch.nn.Sequential(...)
torch.save(model.state_dict(), 'model_weights.pt')
.....
model.load_state_dict(torch.load('model_weights.pt'))
```

Simple Regression problem

$$f(x) = w_0 + w_1 x_i$$

Task : Estimate w_0 and w_1

What is the Loss function ???

Which simple optimizer can be used ??

Hands-on tasks

- Task 1 : Implement Linear regression using Numpy
- Task 2 : Modify task 1 implementation using PyTorch tensors
- Task 3 : Implement linear regression using torch Model
- Task 4 : Add TensorBoard to the implementation
- Task 5 : Use Pretrained models

Resources



PyTorch

<https://pytorch.org>

Thank You